

## Bonus milestones.

You got through the three core milestones. From here, choose one direction and extend your project.

These are intentionally looser than the main milestones. Pick one, combine a few, or write your own. The useful exercise is to decide what PaperClaw should become next, then brief the agent clearly.

---

### Directions worth pushing in

**More formats.** PDFs are the baseline. Add photos with OCR, scanned images, email attachments, or plain `.eml` files. Watch where your classifier starts to break. A badly photographed document is a good test of whether your pipeline is robust or just tuned to the happy path.

**Email ingest.** Pull documents from IMAP or Gmail instead of a local folder. If you connect a real inbox, treat credentials and permissions as part of the design, not an afterthought.

**More tools for the agent.** Search was the first tool. Add another one the agent can call directly:

- *list bills due this week*
- *find all contracts with a given party*
- *surface documents that mention a specific keyword*
- *group receipts by category*

List the next five tools your agent would want, then build one.

**A user-facing interface.** Add a practical entry point for humans: Telegram, voice, a small web UI, or another interface you would actually use. Keep the agent as the worker behind it.

**Document relationships.** Link related documents: a contract, its receipts, the warranty, and the renewal reminder. This lets the agent answer cross-document questions like *"what is the warranty status on that monitor?"*.

**Semantic search.** Add embeddings for the `.md` transcripts. This makes fuzzy queries more useful, for example *"that heating bill from last winter"* or *"the letter about my tax refund"*.

**Auto-action.** Move from answering questions to preparing actions: *"draft the reply", "prepare a payment reminder", "archive everything older than a year"*. Be careful here. Once the tool can act, permissions, confirmation steps, and audit logs matter much more.

**Multi-user.** Support family or small-business use. Different inboxes, a shared library, per-user permissions, and maybe separate agents for different roles.

**Your own idea.** The ideas above are prompts, not a backlog. If you have a better next step, write it down in one or two clear sentences and have the agent build from that brief.

---

## Reminders

The same principles still apply: **decisions, not code** and **feedback loops, not one-shots**.

The foundation you laid in M1 should support the next feature. If it does not, that is useful information. Adjust the design, strengthen the feedback loops, and continue from there.

---

*bonus-milestones.md - optional follow-up work.*