

Milestone 1 - Design & setup.

Goal

Decide what you are building, then set up the repo so an agent can work on it without guessing.

Time estimate

30 minutes (timebox - hard cap)

What to build

- A **design doc** for your version of PaperClaw.
- A **repo** with the basic project structure in place.
- A **library layout** that stores each PDF next to a `.md` **transcript** of its contents.

Your design doc should include

- what your project does
- your **programming language** and key tools / components
- a simple **architecture overview**
- **security and privacy considerations**

The point is to make the important decisions explicit enough that an agent can make sensible implementation choices from them.

Wire up the feedback loops

Before you start building features, add the checks the agent will use to stay on track:

- **Linters and formatters** for your language.
- **Pre-commit hooks** that run automatically.
- **Test scaffolding**. Even one passing test is enough for now.
- A **typecheck** step, if your language has one.

These tools are part of the product. They turn vague review comments into concrete failures an agent can fix.

Acceptance criteria

- A design doc exists, naming stack, architecture, and security/privacy.
- The repo is ready for implementation with an agent: linter, formatter, pre-commit hooks, and test scaffolding are in place.
- The directory layout supports storing PDFs alongside their `.md` transcripts.

What “done” looks like

- Someone can read your design doc and understand what you are building, how you plan to build it, and where the risks are.
- A commit automatically runs the checks you chose.
- You are ready to hand the next task to an agent with a clear brief.