

Milestone 2 - Core inbox workflow.

Goal

Build the core flow: a PDF arrives in the inbox, the system analyzes it, and the document lands in the right place in your library.

Time estimate

45 minutes (recommended)

What to build

- An `inbox/` folder as an input.
- Support for **PDFs** as the core file type.
- Logic that analyzes each file and places it into the `library/` tree with a suitable filename.
- A `.md` **transcript** of the PDF stored next to the original. This carries over from M1's layout and is what lets agents search and read the document later.
- A **log** of what happened during processing.

Implementation notes

- This can be **CLI / file-system based**. No full UI is required.
- Use the library structure you designed in M1. If you change it, update the design doc so the agent has the new context.
- Keep the first version boring. A reliable classifier and file mover is more useful than a nice interface over a fragile core flow.
- If you use an API call for classification or transcription, think about how you will test that path without spending tokens every time.

Acceptance criteria

- A PDF can be added to the `inbox/` folder.
- The system analyzes it.

- The file is moved into the appropriate `library/` subfolder, with a sensible name.
- A `.md` transcript is generated and stored next to the PDF.
- A log is produced, even if it is plain text or markdown.

Strongly recommended

- Add **basic tests**, ideally with mocked API calls where possible. This is not required for the milestone, but it will make the next iteration much easier.

What “done” looks like

- You can demonstrate the full path from incoming PDF to organized, transcribed, logged output.
 - You can explain what happens when the classifier is unsure.
-

Milestone 2 passes if

- a PDF enters the system
- the system processes it
- it ends up in the appropriate `library/` subfolder with a `.md` transcript alongside
- the operation is logged